

- CS324 Final Project Report
 - [Part 0. Project Overview](#)
 - [Part 1. System Setup](#)
 - [Hardware Drive and OS](#)
 - [Software Environments and Packages](#)
 - [Part 2. Network for Person Detection](#)
 - [YOLOv5 Network Architecture](#)
 - [Analysis: Advantages](#)
 - [Part 3. Model Training](#)
 - [Part 4. Person Detection on Jetson Nano](#)

CS324 Final Project Report

陈桂凡 11910503 葛兆宁 11911609 张庭境 11911919 刘仁杰 11911808

Part 0. Project Overview

In this project, we successfully build and run YOLOv5 model inference on Jetson Nano Developer Kit. Results show that our trained model can successfully detect a person if someone appears in front of the camera. We display the detection result (raw image and target detection box in blue) on board and achieve a rate of about 2 to 3 fps for online image streams.

Part 1. System Setup

Hardware Drive and OS

First, we download the [Jetson Nano Developer Kit SD Card Image](#) and write the image to the microSD card by command line instructions using MacOS, according to instructions in [Jetson Nano Developer Kit Start Page](#).

Then, we plug the microSD card into the developer kit box and set up Ubuntu 18.04 when first booting.

Software Environments and Packages

For the first attempt, we install conda-forge which is a community version of anaconda, create a new environment with python 3.9 and install necessary packages using `requirements.txt` provided by [YOLOv5](#), with core packages such as opencv4.1.1, torch1.10.2 and matplotlib3.2.2. However, opencv shows green screen when capturing the images. We find out that the error is incurred by version conflict of gstreamer which is both installed in conda environment and system apt.

Then, we use back the system default python 3.6 with gstreamer in system apt packages and also install the same packages as requirements in python 3.6 enviroment. This time the opencv works normally and the environment is ready for model inference using pytorch.

Part 2. Network for Person Detection

In this project, we choose YOLOv5 as the network for person detection. We look into the network structure in detail in fully understand the algorithm workflow, and alter the input output channels to adapt to inference on Jetson Nano Developer Kit.

YOLOv5 Network Architecture

The diagram illustrates the PANet architecture, divided into three main sections: Backbone, PANet, and Output.

- Backbone:** This section processes the input image through a series of BottleneckCSP blocks. The top BottleneckCSP block is connected to the PANet section. The middle BottleneckCSP block is connected to the PANet section. The bottom BottleneckCSP block is connected to the PANet section.
- PANet:** This section consists of three parallel processing paths. Each path starts with a BottleneckCSP block, followed by a Conv1x1 layer, an UpSample layer, and a Concat layer. The Concat layers are connected to the corresponding BottleneckCSP blocks in the Output section.
- Output:** This section takes the output from the PANet section and processes it through a series of Conv1x1 layers. The top Conv1x1 layer is connected to the PANet section. The middle Conv1x1 layer is connected to the PANet section. The bottom Conv1x1 layer is connected to the PANet section.

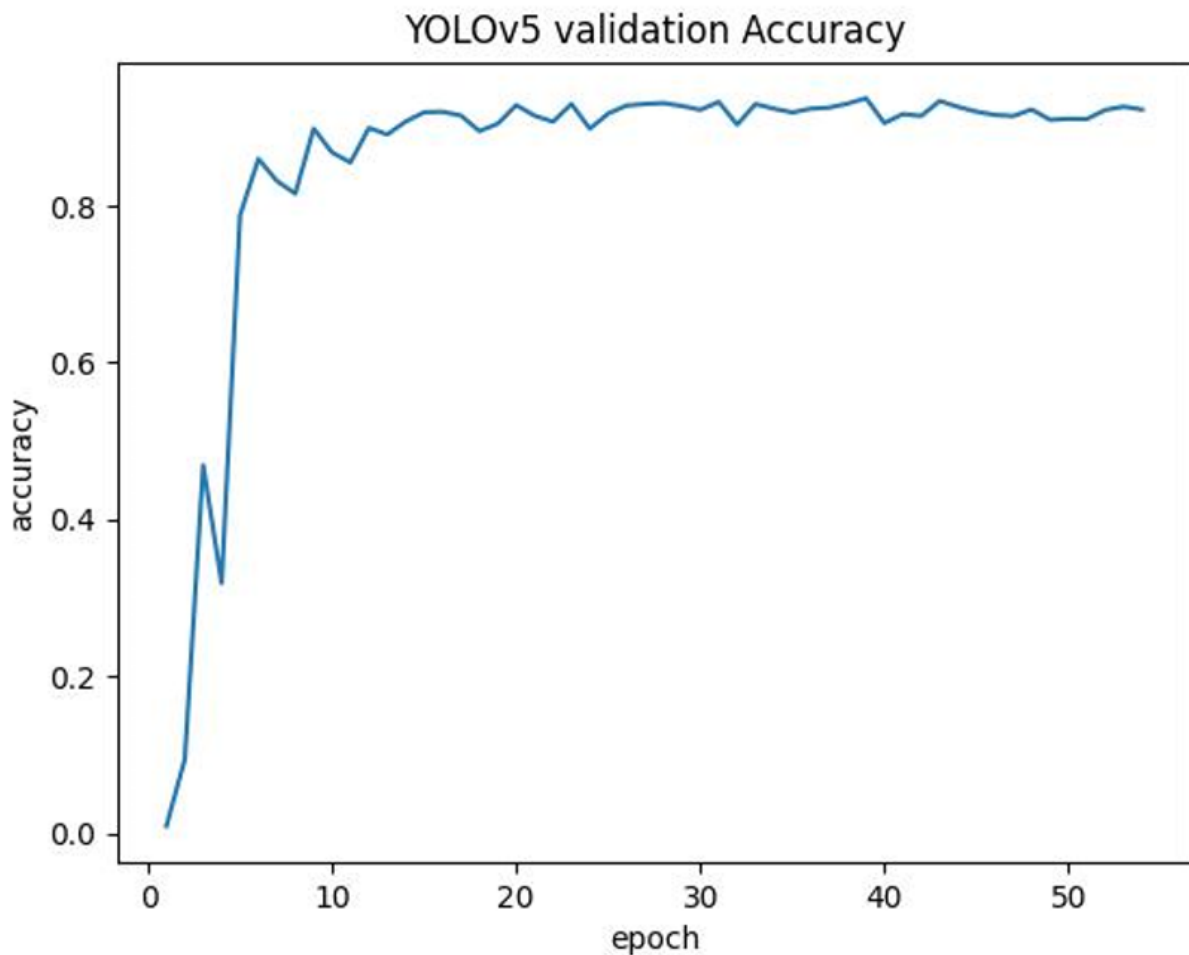
Analysis: Advantages

Part 3. Model Training

The full name of IoU is Intersection over Union, which refers to the intersection area between the predicted content and the actual content divided by the sum of their areas. The full name of mAP is mean average precision. It means to evaluate the precision of

each batch of pictures under test when the IoU is determined, and then take the average value of this batch of pictures. mAP_0.5 is a suitable measurement index for this project, which refers to the mAP when the IoU is at 50%.

The results demonstrates that the YOLOv5 learns well on the small datasets with 1500 images, which achieves an accuracy over 90%.



Part 4. Person Detection on Jetson Nano

The procedure of person detection on Jetson nano can be roughly divided into three steps:

- We use opencv to fetch online image stream captured from the camera on board.
- For each frame of image stream, we use the trained YOLOv5 to detect whether there is a person.
- Display the detection box and original image together on the external display.

Results shows that we can successfully detect a person with a detection rate of 2 to 3 fps:



In tests, people wearing masks or upside down can also be detected.